



Speech Therapy Suite

Technical Reference Documentation

Version 4.0.2 (May 2026)

A customisable, browser-based speech-language practice suite.
Runs as a Claude.ai artifact or locally via Vite. Powered by Claude Sonnet.

1. Architecture Overview

The Speech Therapy Suite is a single-page React application that runs entirely in the browser. There is no backend server — all persistence is via localStorage and all AI calls are made directly from the browser to the Anthropic API.

1.1 Technology Stack

Layer	Technology
UI framework	React 18 (hooks only, no class components)
Build tool	Vite 5 with ESM modules
AI / LLM	Anthropic Claude (claude-sonnet-4-5 or latest)
Styling	Inline styles only — no CSS framework
Speech synthesis	meSpeak v2 (eSpeak compiled to JavaScript via Emscripten)
Distribution	Claude.ai artifact (single JSX file) or Vite dev/build
Persistence	localStorage (all keys prefixed ppa_)
PWA	manifest.json + service worker for iPad home-screen install

1.2 Source File Layout

```
ppa-source/  
  ppa-speech-therapy_main.jsx  # App shell + all modules except Naming and SentenceBuilder  
  NamingModule.jsx            # Picture-naming practice with spaced repetition  
  SentenceBuilderModule.jsx    # Visual drag-and-drop sentence construction  
  ExportImportSystem.jsx      # .ppa / .ppabak export, import, and backup logic  
  shared.jsx                  # Shared utilities: fetchAnthropicApi, CallAPI, ThinkingDots, ZoomableGraphic  
  data/  
    config.js                 # CLAUDE_MODEL constant + Dr. Aria SYSTEM_PROMPT  
    dictionary.js             # Unified word->{graphic, cues, categories} store  
    namingItems.js            # 10 built-in picture-naming items (seed data)  
    sbWordBank.js             # Noun/verb/adjective/adverb/pronoun/prep/article banks  
    sbConjugation.js          # Verb conjugation rules for all tenses  
    repetitionItems.js        # Repetition drill levels  
    sentenceTasks.js          # Sentence completion and construction prompts  
    scripts.js                # Functional phrase scripts  
    assessmentTasks.js        # Evaluation items  
    videoClips.js             # Video comprehension clips and question types  
    tools.js                  # Sidebar navigation definitions
```

1.3 Module Dependencies

All inter-module shared code lives in shared.jsx. Modules import only what they need. The dependency graph is intentionally flat — no circular imports.

2. Shared Utilities (shared.jsx)

All Anthropic API access and the loading indicator live in `shared.jsx`. Never duplicate these inline in a module.

2.1 `fetchAnthropicApi(body, signal?)`

Low-level async helper. Applies all required headers (`x-api-key`, `anthropic-version`, `anthropic-dangerous-direct-browser-access`). Returns parsed JSON. Throws on network error or abort. Use this for fire-and-forget calls (e.g. emoji lookup in `SentenceBuilder`).

```
const data = await fetchAnthropicApi({
  model: CLAUDE_MODEL,
  max_tokens: 256,
  messages: [{ role: 'user', content: prompt }]
}, abortSignal);
```

2.2

React component that fires one API request on mount and calls `onResult(text)` or `onError(err)` exactly once. Uses `AbortController` — the request is cancelled automatically on unmount. `onResult` always receives a non-empty string (falls back to "Well done — keep going!"). Mount it conditionally: `{pendingAI && <CallAPI ... />}`.

2.3

Animated three-dot spinner. Use during any AI loading state.

2.4

Wraps any graphic (emoji or image) and makes it clickable. On click, a full-screen modal overlay opens. Inside the overlay:

- **+ / –** buttons (top-right corner) increment or decrement `zoomScale` by 0.25 (minimum 0.5).
- **Corner drag handle** (bottom-right, L-shaped border, nwse-resize cursor) — pointer events on the handle update `zoomScale` continuously as the pointer moves: `scale += (dx + dy) / 200`.
- Image display size is `width * 2 * zoomScale × height * 2 * zoomScale`, capped at 90 vw / 90 vh.
- Pressing **Escape** or clicking outside the card closes the overlay and resets `zoomScale` to 1.

Props: `graphic` (string — emoji or image `src`), `alt` (string), `width` / `height` (px, for image elements), `fontSize` (px, for emoji spans), `imgStyle` / `spanStyle` (merged into the display element).

Used by: `NamingModule` (practice card, item list) and `AssessmentModule` (confrontation naming task). Import from `shared.jsx`; never inline.

3. Data Layer

3.1 localStorage Keys

Key	Contents
ppa_dictionary	Canonical word→{graphic, cues, categories} map (JSON)
ppa_naming_items	Current naming practice list (JSON array)
ppa_naming_sr	Spaced repetition state per word (JSON object)
ppa_naming_audio_hints	Boolean — audio hints on/off (string 'true'/'false')
ppa_session_log	Array of session log entries (JSON)
ppa_progress	Accuracy history per module (JSON)
ppa_scripts	Custom scripts (JSON array)
ppa_sentences	Custom sentence tasks (JSON array)
ppa_video_clips	Custom video clips (JSON array)
ppa_assessment	Custom assessment tasks (JSON array)
fam_persons	Family graph — array of person objects {id, name, photo_url, facts}
fam_relationships	Family graph — array of edge objects {id, from, to, type, lineage}
fam_profile	Patient profile {name, conditionType, ...} — shared with ProfileModule

3.2 Dictionary (data/dictionary.js)

Single source of truth for all word graphics, stored in localStorage under ppa_dictionary. Both NamingModule and SentenceBuilder read and write through the dictionary API:

- dictLoadNamingItems() / dictSaveNamingItems(items) — load/persist the naming practice list.
- dictGetGraphic(word, fallback) — resolve canonical emoji or base64 image.
- dictAddWord(word, graphic) — register a new word (first writer wins; ■ is always upgradeable).
- useDictionaryLookup() — React hook returning a stable { word: graphic } map.

4. Export / Import System (ExportImportSystem.jsx)

4.1 File Formats

.ppa files — per-module item exports. Format: { ppaExport: true, moduleId, items: [...] }.

.ppabak files — full-app backup of all ppa_* localStorage keys. Restoring reloads the page.

4.2 Public API

All public functions and components are named exports. The main app and each module import only what they need.

Export	Type	Purpose
ppaDownload(filename, data)	function	Triggers browser file download
ppaDoBackup()	function	Exports full .ppabak backup
ppaBackupsIsStale()	function	Returns true if last backup >7 days ago
PpaAdminToolbar	component	Admin mode toolbar strip
PpaExportDialog	component	Export dialog for a module
PpaReexportDialog	component	Re-export changed items dialog
BackupRestorePanel	component	Full backup/restore UI panel

5. Naming Module (NamingModule.jsx)

5.1 Practice Phases

The naming flow moves through these phases for each item:

Phase	Display	Transitions
show	Graphic + empty input	Space → phoneme hint; Got it! or Enter → record correct; Semantic Cue btn → semantic
semantic	Semantic cue text shown	Got it! or Enter → record semantic_cued; Phonemic Cue btn → phonemic
phonemic	Phonemic cue shown	Got it! or Enter → record phonemic_cued; Reveal btn → answer
answer	Full word shown	Next word → advance

5.2 Spaced Repetition Engine

PPA-adapted SR — deliberately conservative (max 5-day interval, regression decay on every load):

Result	SR_FACTOR	When used
correct	×1.2	Answered without any help
space_cued	×0.9	Used the Space-key phoneme starter
semantic_cued	×0.8	Used the concept hint
phonemic_cued	×0.5	Used the sound cue
failed	reset → 1 day	Needed full reveal

SR state is stored in localStorage under ppa_naming_sr: { [word]: { interval, dueDate, streak, lastResult, lastSeen } }.

Words answered with phonemic_cued or failed are re-inserted 4 positions ahead in the session queue for same-session repetition.

5.3 AI Feedback Flow

After recording any response, <CallAPI> is mounted to fetch Dr. Aria's feedback. Once the response arrives the "Next word →" button appears. If the API call fails (no key, CORS, etc.) onError calls next() directly so the user is never left stuck.

5.4 Phoneme Starter (Space Key) with Audio Hints

In the show phase, when the response input is empty, pressing **Space** reveals a prefix of the target word and (if audio hints are on) speaks it aloud via **meSpeak**.

Audio-hints toggle

A  button in the bottom-right corner of the input area toggles audio hints on/off. The choice is persisted in localStorage under ppa_naming_audio_hints. Default is on.

Audio-on mode (default)

Each Space press advances the display to the next vowel boundary (via `findNextVowelEnd`) plus the immediately following consonant group (via `getNextPhoneme` as an *anchor*), then speaks that exact text:

```
const newCharCount = findNextVowelEnd(item.word, phonemesRevealed);
const anchor = newCharCount < item.word.length
  ? getNextPhoneme(item.word, newCharCount) // next consonant group
  : "";
const displayCount = newCharCount + anchor.length;
setPhonemesRevealed(displayCount);
meSpeak.resetQueue();
meSpeak.speak(item.word.slice(0, displayCount).toLowerCase(), { speed: 130 });
```

The trailing consonant anchor is essential for correct vowel quality. Without it, eSpeak's LTS rules assign the wrong allophone to word-final vowels: `meSpeak.speak("gla")` produces a schwa-like sound. With the anchor, `meSpeak.speak("glas")` forces a CVC (closed syllable) context and correctly produces the short /æ/ of *glasses*. Similarly, speaking "chair" gives eSpeak's dictionary pronunciation /tʃaɪr/ rather than the standalone-word pronunciation /tʃaɪ/ of "chai".

Display and audio always show the same characters — the patient never hears more than is displayed.

Example	Display	speakText	Result
glasses, press 1	GLAS	"glas"	short-/æ/ ✓
glasses, press 2	GLASSES	"glasses"	full word ✓
chair, press 1	CHAIR	"chair"	/tʃaɪr/ not /tʃaɪ/ ✓
apple, press 1	AP	"ap"	short-/æ/ ✓

Audio-off mode

Each Space press advances by one grapheme group (1–3 characters) using `getNextPhoneme`. No audio is played.

Vowel and grapheme helpers

Function	Purpose
<code>findNextVowelEnd(word, startPos)</code>	Returns char position after next vowel group (single or digraph).
<code>getNextPhoneme(word, pos)</code>	Returns next grapheme cluster (1–3 chars) via TRIGRAPHS / DIGRAPHS arrays; used for both audio-off advances and the audio-on anchor.
VOWEL_GROUPS	Multi-letter vowel digraphs (ai, ay, ea, ee, igh, ough, ...) checked longest-first.

Function	Purpose
LETTER_VOWELS	Set: {a, e, i, o, u}.

meSpeak initialisation

```
import meSpeak from "mespeak";
import meSpeakConfig from "mespeak/src/mespeak_config.json";
import enVoice from "mespeak/voices/en/en-us.json";

if (!meSpeak.isConfigLoaded()) meSpeak.loadConfig(meSpeakConfig);
if (!meSpeak.isVoiceLoaded()) meSpeak.loadVoice(enVoice);
```

Initialisation runs at module level. `meSpeak.resetQueue()` is called before each `speak()` to prevent audio queuing when Space is pressed rapidly.

SR result

Submitting with `phonemesRevealed > 0` records `space_cued` (`SR_FACTOR × 0.9`). The mild penalty acknowledges the word was not recalled unaided, while reflecting that audio support is lighter than semantic or phonemic cueing.

5.5 AI-Assisted Item Defaults

When the **Word** field in the Add Item form (AdminPanel → ItemForm) loses focus and contains a non-empty value, generateDefaults(word) is called via fetchAnthropicApi. The function requests a JSON object from the Claude API containing four fields:

- **graphic** — a single emoji that best represents the word
- **category** — a short category label (e.g. "fruit")
- **clue_semantic** — a short sentence cue (e.g. "It's a fruit you eat")
- **clue_phonemic** — a phonemic cue noting the starting sound (e.g. "Starts with 'A'...")

Each field is applied only if it is still at its blank default (■ for graphic, empty string for the rest) — any value already typed by the admin is preserved. A <ThinkingDots /> spinner is shown during the request. Failures are silent; all fields remain editable. Requires an internet connection and a valid API key.

5.6 BulkImportPanel

Renders inside AdminPanel when mode === "bulkimport". Lets the admin select multiple image files at once (drag-and-drop or multi-select file browser; accepts image/*, .heic, .heif). Each file is compressed to a base64 JPEG via compressImageFile (240 px max side, 0.75 quality) and added as a draft row in an editable grid table.

Grid columns

- **Photo** — compressed thumbnail (52 × 52 px).
- **Word** — text input; blurring the field fires generateForDraft(id, word), which calls fetchAnthropicApi to suggest category, clue_semantic, and clue_phonemic. Only empty fields are filled. A ThinkingDots spinner appears during the request.
- **Category, Semantic cue, Phonemic cue** — free-text inputs.
- **Remove** — deletes the row from the draft list.

"Generate AI defaults for all" iterates the draft list sequentially, calling generateForDraft for each row that has a non-empty word.

"Save N items" calls onSaveAll passing rows where all four fields (word, category, clue_semantic, clue_phonemic) are non-empty. handleBulkSave deduplicates against the existing item list by word (case-insensitive) before appending.

Filename inference: camera-roll names matching /^(img|dsc|photo|picture|p|image|pic)[-_]?[a-z0-9]+\$/i are left blank; other filenames are lowercased with hyphens and underscores converted to spaces and used as the initial word value.

5.7 Generate by Category

Renders inside AdminPanel when mode === "generate". Lets the admin auto-populate a category with the most common nouns using a single AI call.

Controls

- **Category input** — free-text field for the category name (e.g. "fruit", "animal", "clothing").
- **Count** — number input, 1–50, default 15.
- **Generate** — fires a single fetchAnthropicApi call requesting the N most common nouns in the category. The prompt asks for a JSON array of objects, each with word, graphic (emoji), clue_semantic, and clue_phonemic. A <ThinkingDots /> spinner is shown while the request is in flight.

Results grid

Uses the same grid style as BulkImportPanel. Columns: checkbox, Graphic, Word, Category, Semantic cue, Phonemic cue, Status.

Each generated noun is classified via checkDuplicate():

- **New items** — shown with a green **NEW** badge and pre-selected checkbox. Category is set to the new category name.
- **Duplicates** — shown with an amber **EXISTS** badge, no checkbox, and all existing values including the original category displayed in amber. Duplicates are never added.

A summary bar above the grid shows the count of new and existing items, plus a "Select all new" checkbox. The grid scrolls independently; controls and the "Add N items to library" action button remain pinned.

Actions

"Add N items to library" calls handleAddGenerated(), which filters to selected new items, strips the internal _status field, assigns a custom-{timestamp} id to each, appends them to the item list via onUpdate, and returns to list view.

If the API call fails (no key, network error), an alert is shown and the controls remain active so the admin can retry.

6. Sentence Builder Module (SentenceBuilderModule.jsx)

6.1 Word Bank

Words are organised into seven categories: nouns, verbs, adjectives, adverbs, pronouns, prepositions, and articles. The bank is stored in `data/sbWordBank.js`. Conjugation rules for all tenses are in `data/sbConjugation.js`.

6.2 Drag-and-Drop

Built with native HTML5 drag-and-drop events (no third-party library). Touch events on iOS are handled with a polyfill embedded in the component. Tiles can be dragged from the bank into the tray, and reordered within the tray.

6.3 AI Feedback

When the patient taps **Check sentence**, the constructed sentence is sent to `fetchAnthropicApi` with a prompt asking Dr. Aria to give brief, encouraging grammatical feedback. The emoji for new words is also looked up via `fetchAnthropicApi` and stored in the dictionary.

7. Other Modules

7.1 AI Therapist (TherapistModule)

Full multi-turn conversation with Dr. Aria. Messages are accumulated in component state (not persisted). Uses <CallAPI> mounted on every send. The system prompt is defined in data/config.js.

7.2 Assessment

Presents tasks from data/assessmentTasks.js. Responses are recorded as correct/incorrect/partial and logged to the session log. Admin can add, edit, and reorder tasks.

7.3 Repetition

Drills from data/repetitionItems.js, grouped by level. Accuracy is tracked per level. Dr. Aria gives end-of-level feedback.

7.4 Script Training

Scripts from data/scripts.js. Each script has a name, context, and ordered lines. After completing a script, Dr. Aria gives feedback. Scripts can be exported/imported as .ppa files.

7.5 Sentence Work

Two sub-modes (completion and construction) drawing from data/sentenceTasks.js. Both use <CallAPI> for per-item AI feedback.

7.6 Video Questions

Clips defined in data/videoClips.js. YouTube embeds via iframe. Questions can be multiple-choice (EMOJI_OPTIONS) or open-ended. Responses are logged.

8. Family Module (FamilyModule.jsx)

8.1 Graph Data Model

The family data is stored as a property graph in two localStorage keys: `fam_persons` and `fam_relationships`. Only *direct* relationship types are stored as edges — parent, partner, and sibling. All other labels (grandparent, grandchild, aunt/uncle, cousin, etc.) are derived at render time via BFS traversal.

Stored type	Edge direction	Lineage variants
parent	from: parent, to: child	biological · step · adopted · guardian · half
partner	symmetric (from/to both stored)	biological (default)
sibling	symmetric (from/to both stored)	biological · half

8.2 BFS Relationship Derivation

`deriveLabel(personId, relationships)` runs a BFS from the special `_patient` node to `personId` and collects a path of directional steps. `pathToLabel(path)` maps the step sequence to a human-readable string:

Path pattern	Label
up	parent
down	child
side (partner)	partner
side (sibling)	sibling
up, up	grandparent
down, down	grandchild
up, down	sibling (via shared parent)
up, up, down	aunt / uncle
up, down, down	niece / nephew
up, up, up	great-grandparent
up, up, down, down	cousin

8.3 Tree Layout

`buildTreeLayout(persons, relationships)` assigns generation numbers via BFS from `_patient` (parent edges = ± 1 , partner/sibling = 0), then positions each generation as a horizontal row. The patient node is always at the horizontal centre of the SVG. Gen-0 (same generation as patient) is ordered: siblings | patient | partners | others.

`buildLines(persons, relationships, positions)` draws three line types:

- **Teal pedigree bars** — children sharing the same parent set share a single horizontal bar with vertical drops (no false sibling connections).

- **Rose double-lines** (stroke #B05090) — one line at $y-3$ and $y+3$ for each partner pair; drawn once per pair via a drawnPairs Set.
- **Dashed teal lines** — siblings only when no shared parent is already in the graph (avoids duplication with bar lines).

8.4 Legacy Migration

On first load, loadFamilyMembers() checks for the old fam_family_members key. If found, migrateLegacy() converts each record: parent → parent edge to patient; child → parent edge from patient; partner / sibling → symmetric edges. Derived types (grandparent, cousin, etc.) become persons with no edges — the user must add the intermediate links manually. The legacy key is then removed.

9. Progress Module

Reads from the session log (ppa_session_log) and the SR state (ppa_naming_sr) to display:

- **Session log** — timestamped list of all activity this session.
- **Accuracy charts** — per-module correct/incorrect percentages.
- **SR table** — per-word interval, streak, last result, next due date.

The session log is not persisted across page reloads — it is held in React state at the app root and passed down as a prop.

10. Configuration (data/config.js)

config.js exports two constants:

CLAUDE_MODEL — the Anthropic model string used in all API calls (e.g. "claude-sonnet-4-5"). Change this to upgrade to a newer model.

SYSTEM_PROMPT — Dr. Aria's full system prompt. Defines her persona, tone, and therapeutic approach. Edit this to customise her behaviour.

10.1 Environment Variable

The Anthropic API key is read from `import.meta.env.VITE_ANTHROPIC_API_KEY`. Create a `.env` file in the repo root:

```
VITE_ANTHROPIC_API_KEY=sk-ant-...
```

Without it the AI calls fail silently and the app auto-advances rather than showing Dr. Aria's feedback.

11. Implementation Notes

11.1 React StrictMode

StrictMode is active in development (src/main.jsx). Effects run twice. All useEffect hooks must return a cleanup function. <CallAPI>'s AbortController handles this automatically.

11.2 No Backend — Security Considerations

The API key is stored in browser localStorage (or .env for Vite builds). This is acceptable for a single-patient, single-device use case but is not suitable for a multi-user web deployment. In that scenario the key must be moved to a backend proxy.

11.3 CORS — anthropic-dangerous-direct-browser-access

Anthropic's API requires the header anthropic-dangerous-direct-browser-access: true for direct browser calls. This is set in fetchAnthropicApi and must not be removed.

11.4 Admin PIN

The PIN is defined as a constant in NamingModule.jsx (const ADMIN_PIN = "1234"). Change this before clinical deployment. The PIN is checked only in the browser — it provides no server-side security.

11.5 PWA / Service Worker

The pwa/ directory contains a manifest.json and a service-worker.js that cache the app shell for offline use. The service worker is registered in index.html. AI features remain unavailable offline.

11.6 localStorage Limits

Most browsers allow 5–10 MB of localStorage per origin. Base64-encoded photos in the dictionary are the largest consumers. The app does not enforce a limit — if storage fills, writes will silently fail. Advise users to use emoji graphics rather than photos where possible.

11.7 Bundle Size

The Vite build produces a single JS bundle. The meSpeak library (eSpeak WASM) adds ~1.5 MB to the bundle. This is acceptable for a local/LAN deployment but may be slow on a first load over a slow connection. Consider using import() lazy loading for meSpeak if load time is a concern.

11.8 Audio Hints — Vowel Anchoring Approach

The initial approach used phonemizer (Xenova, eSpeak-NG WASM) to convert each word to IPA, map IPA tokens to eSpeak X-SAMPA notation, and call meSpeak.speak("[[g l {}]]") in phoneme mode. Two problems made this unworkable: (1) the phonemizer package failed to initialise its search engine in the Vite development environment; (2) meSpeak v2's [[...]] phoneme notation produced no audio output (returns null for rawdata: true).

The replacement approach requires no external phonemizer. Speaking the prefix plus one anchor consonant in text mode (meSpeak.speak("glas")) achieves correct vowel quality because eSpeak's LTS rules handle closed-syllable /æ/ correctly. For dictionary words the full matched string invokes eSpeak's internal dictionary, giving the correct pronunciation automatically.

phonemizer and espeak-ng have been removed from package.json. Both are listed in vite.config.js → optimizeDeps.exclude to prevent Vite from attempting to pre-bundle any stale package stubs.

12. Installer / Distribution

12.1 macOS Installer (installer/install.sh)

A shell script that:

- Checks for Node.js and npm.
- Runs npm install.
- Copies PDF documentation to the install directory.
- Performs a preflight check for PDF presence.
- Launches the Vite dev server.

12.2 Windows Installer (win-installer/install.ps1)

A PowerShell script performing the same steps as the macOS installer, adapted for Windows paths and conventions.

12.3 Claude.ai Artifact

The bundle files (ppa-speech-therapy-bundle.jsx) in each installer directory are single-file versions of the app suitable for pasting into a Claude.ai artifact. They include all module code inlined.

13. Changelog

Version 4.0.2 (May 2026)

- FamilyModule: replaced flat relation_type with a proper graph model (fam_persons + fam_relationships). Direct edges only (parent / partner / sibling); all other labels derived by BFS at render time.
- Family tree SVG: patient centred horizontally; teal pedigree bars for parent–child groups; rose double-lines for partners; dashed teal for siblings with no recorded shared parent.
- PersonForm: multi-link editing with sticky "related to" default; automatic migration from legacy fam_family_members key.
- NamingModule: Enter / Return in the answer input fires Got it! across all three practice phases.
- NamingModule: practice buttons gain a visible focus ring for keyboard tab-navigation.
- ZoomableGraphic: +/- buttons and bottom-right corner drag handle for free-form resize in the zoom overlay.

Version 4.0.1 (March 2026)

- Added click-to-zoom for graphics: tapping any picture card in Naming Practice or Assessment opens a 2× popup. Press Escape or click outside to close.
- Added ZoomableGraphic shared component to shared.jsx.

Version 4.0.0 (March 2026)

- Added meSpeak audio hints with vowel-anchored pronunciation in Naming module.
- Replaced phonemizer / espeak-ng IPA pipeline with text-mode meSpeak.
- Added ppa_naming_audio_hints localStorage key for toggle persistence.
- Added Space-bar phoneme starter (space_cued, SR ×0.9).
- Extracted shared utilities to shared.jsx (fetchAnthropicApi, CallAPI, ThinkingDots).
- Renamed main source file to ppa-speech-therapy_main.jsx.
- Added Vite project scaffold (package.json, vite.config.js, index.html, src/main.jsx).
- Added iPad PWA (manifest.json + service worker).
- PDF documentation bundled into macOS and Windows installers.
- Max SR interval corrected to 5 days.